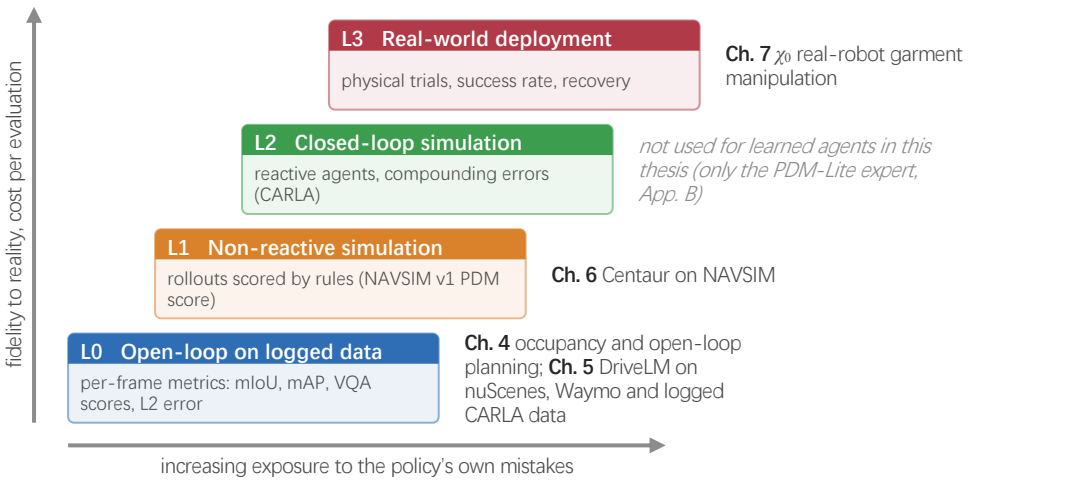
fidelity to reality, cost per evaluation
L3 Real-world deployment
physical trials, success rate, recovery
Ch. 7 χ_0 real-robot garment manipulation
L2 Closed-loop simulation
reactive agents, compounding errors (CARLA)
not used for learned agents in this thesis (only the PDM-Lite expert, App. B)
L1 Non-reactive simulation
rollouts scored by rules (NAVSIM v1 PDM score)
Ch. 6 Centaur on NAVSIM
L0 Open-loop on logged data
per-frame metrics: mIoU, mAP, VQA scores, L2 error
Ch. 4 occupancy and open-loop planning; Ch. 5 DriveLM on nuScenes, Waymo and logged CARLA data
increasing exposure to the policy's own mistakes